

Programming Assignment 4: Naïve Bayes Text Classification

Assigned: April 14

Due: May 5

In this assignment you will use the Naïve Bayes method to learn a classifier for text from training data, and use a test set to evaluate the quality of the classifier.

1 Input format

The major input to the program consists of a file of short biographies (modified from Knowledge Graph and Wikipedia) of various people, tagged with a category. There are two samples available: a small one at <http://www.cs.nyu.edu/faculty/davise/ai/bioCorpus.txt> and an tiny one at <http://www.cs.nyu.edu/faculty/davise/ai/tinyCorpus.txt>

The format of this file is as follows:

Separate biographies are separated by 1 or more blank lines.

In each biography, the first line is the name of the person. The second line is the category (a single word). The remaining lines are the biography. You may assume that the characters in the file are all a-z, A-Z, comma, period, and white space; there are no numbers and no other punctuation.

This should be read in from a file “corpus.txt” in the same working directory as the executable program.

The second input is a user input of an integer N, which is the number of entries in the corpus to use as the training set. The program will use the first N biographies in the corpus as the training set and the rest as the test set. This should be read in from standard input or in some other convenient way from the user. (If you do something other than standard input, please alert the grader in a ReadMe file.)

A third, optional, input is a file “stopwords.txt”, containing a list of stop words. This should also be in the same working directory. This is available at <http://www.cs.nyu.edu/faculty/davise/ai/stopwords.txt>. It is optional, because it is OK just to hard-code the list of stop words, if you prefer.

2 Output format

The output contains:

- For each person in the test set, the probabilities associated with each category; the prediction (the category with the highest probability), and a statement whether this is right or wrong.
- The overall accuracy of the classifier on the test set.

For instance, for the tiny corpus linked above, with $N=5$, the output is:

Benjamin Disraeli. Prediction: Writer. Wrong.
Government: 0.44 Music: 0.07 Writer: 0.48

George Eliot. Prediction: Writer. Right.
Government: 0.01 Music: 0.07 Writer: 0.91

Barbara Jordan. Prediction: Government. Right.
Government: 0.97. Music: 0.03. Writer: 0.00

Clara Schumann. Prediction: Music. Right
Government: 0.01. Music: 0.98. Writer: 0.01

Overall accuracy: 3 out of 4 = 0.75.

Probabilities should be given to an accuracy of 0.01, as above.

This output is available at <http://www.cs.nyu.edu/faculty/davise/ai/tinyOutput.txt>.

I will use this as an example throughout the assignment. The “tiny training corpus” is the first 5 biographies in the tiny corpus, and the “tiny test corpus” is the last 4.

3 The Program

The program is divided into two parts: Learning the classifier from the training set (the first N biographies in the input), and evaluating the classifier over the test set (the remaining biographies).

3.1 Learning the classifier from the training set

3.1.1 Normalization

In reading in the training set:

- All words in the biography should be normalized to lower case.
- Stop words should be omitted. Stop words are: (a) any word of one or two letters; (b) any word in the list of stopwords provided. <http://www.cs.nyu.edu/faculty/davise/ai/stopwords.txt>.

3.1.2 Counting

Compile a count of

- For each category C , the number of biographies of category C in the training set T , $\text{Occ}_T(C)$.
- For each category C and word W in the training set T , the number of biographies of category C that contain word W , $\text{Occ}_T(W|C)$. It does not matter how many times a word occurs within a given biography, as long as it occurs once (this is known as the “Bernoulli method” or, more generally, as a “set of words” model).

Note: THE WORD USED FOR THE CATEGORY DOES NOT COUNT AS A WORD IN THE TEXT. So, for example, in the biography of Bella Abzug, “Government” is not considered a word W in the biography.

3.1.3 Probabilities

Step 1: For each classification C , define $\text{Freq}_T(C) = \text{Occ}_T(C)/|T|$, the fraction of the biographies that are of category C . For instance, in the tiny training corpus, $\text{Freq}_T(\text{Government}) = 2/5$.

For each classification C and word W , define $\text{Freq}_T(W|C) = \text{Occ}_T(W|C)/\text{Occ}_T(C)$, the fraction of biographies of category C that contain W . For instance, in the tiny training corpus, $\text{Freq}_T(\text{“american”}|\text{Government}) = 1/2$, since there are two biographies of category *Government* and one of these has the word “american”.

Step 2: For each classification C and word W , compute the probabilities using the Laplacian correction. Let $\epsilon = 0.1$.

$$P(C) = \frac{\text{Freq}_T(C) + \epsilon}{1 + \text{number of categories} * \epsilon} \quad P(W|C) = \frac{\text{Freq}_T(W|C) + \epsilon}{1 + 2 * \epsilon}$$

In the second equation, the “2” in the denominator reflects the fact that each random variable W has two possible outcomes: true and false.

For instance $P(\text{Government}) = (0.4 + 0.1)/(1 + 3 * 0.1) = 0.3846$.
 $P(\text{“american”}|\text{Government}) = (0.5 + 0.1)/(1 + 2 * 0.1) = 0.5$.

Step 3: Compute negative log probabilities to avoid underflow. The base of the logarithm doesn’t matter, except that you will have to use the same base when you reconstruct the probabilities in step 4.b of part 2. I’ve used base 2 in computing the trace. For each classification C define $L(C) = -\log_2(P(C))$ and define $L(W|C) = -\log_2(P(W|C))$. For instance, $L(\text{Government}) = 1.3785$ and $L(\text{“american”}|\text{Government}) = 1.0$.

Now you’re done with the learning phase.

3.2 Applying the classifier to the test data

For each biography B :

- Step 1. Normalize the text as in part I. Additionally skip any word that did not appear at all in the training data.
- Step 2. For each category C , compute $L(C|B) = L(C) + \sum_{W \in B} L(W|C)$. As in the learning phase:
 - Each word counts only once, no matter how often it appears in the biography.
 - DO NOT USE THE CATEGORY ATTACHED TO THE BIOGRAPHY. This is even more important now than it was in part I. Whenever you use part of a labelled corpus as a test set, you obviously have to ignore the label; otherwise, the learning algorithm is just cheating off a crib sheet.
- Step 3. The prediction of the algorithm is the category C with the smallest value of $L(C|B)$. It is extremely unlikely that you will ever get a tie; if you do, it may be broken arbitrarily (i.e. don’t worry about the case.)

Step 4. To recover the actual probabilities: Let k be the number of categories.

- 4.a For $i = 1 \dots k$ let $c_i = L(C_i|B)$, the values of L for all the different categories. Let $m = \min_i c_i$, the smallest value of these.
- 4.b For $i = 1 \dots k$, if $c_i - m < 7$ then $x_i = 2^{m-c_i}$ else $x_i = 0$. The base 2 here is the same one as used in step 3 of the learning algorithm.
- 4.c Let $s = \sum_i x_i$, For $i = 1 \dots k$, $P(C_i|B) = x_i/s$.

There is a trace of the execution of the algorithm for the tiny corpus with $N = 5$ at <http://www.cs.nyu.edu/faculty/davise/ai/tinyChart.txt>.

4 Program requirements

You may hard code the list of stop words.

You may not hard code the set of categories; e.g., do not assume that the only categories are Government, Music, and Writer.

If you want, you may assume that the entire corpus fits in memory, so it is OK to begin by reading in the whole file corpus.txt and then working on it as an in-memory data structure. However, it is preferable, and actually not very difficult, to write the program so that it works by doing a single pass reading through the file; it just requires carefully keeping track of things. The memory required for doing this the right way is $O(\text{number of different words} \cdot \text{number of different categories})$ which in general is much smaller than the total size of the corpus.

Similarly, you may make reasonable assumptions on the number of different words (not more than 10,000), the number of different categories (not more than 100) the maximum length of a biography (not more than 200 words), and the total number of biographies (not more than 1000) if you want. It is certainly better to write your code so that it doesn't need these kinds of assumptions, however.

In the learning phase, you should not write code whose time requirement is $\Omega(D^2)$ where D is the number of different words. $O(D \log D \cdot k + |T|)$ is acceptable (k is the number of categories; $|T|$ is the total size of the training set), but really you should achieve $O(D \cdot k + |T|)$ expected time (hint: hash tables). In the classification phase, the time to classify each biography B should be $O(|B| \cdot k \cdot \log D)$ or preferably $O(|B| \cdot k)$ where $|B|$ is the length of the biography. A penalty of 5 points out of 100 will be imposed on any code that requires time $\Omega(D^2)$ in the learning phase or $\Omega(D)$ in the classification phase or both (maximum of 5 points total penalty).

You may assume that the input is correctly formatted, as specified in section 1 of the assignment.